

Real-Time and Embedded Systems

Lab 7

Alex Hoffman

`alex.hoffman@tum.de`

TU München
Institute for Real-Time Computer Systems
(RCS)

2020



Lehrstuhl für
Realzeit-Computersysteme



WCET with Caches

Problem:

Due to caches, execution time of individual basic blocks can vary at different points in the program execution.

WCET with Caches

Problem:

Due to caches, execution time of individual basic blocks can vary at different points in the program execution.

What can we do?

WCET with Caches – What can we do?

WCET with Caches – What can we do?

1 Just ignore them

- **no:** This will lead to huge over- or underestimations up to factor of 100 on modern processors.

WCET with Caches – What can we do?

- 1 Just ignore them
 - **no:** This will lead to huge over- or underestimations up to factor of 100 on modern processors.
- 2 Find longest path with ILP technique, and then simulate it to get cache misses/hits

WCET with Caches – What can we do?

- 1 Just ignore them
 - **no:** This will lead to huge over- or underestimations up to factor of 100 on modern processors.
- 2 Find longest path with ILP technique, and then simulate it to get cache misses/hits
 - **no:** Longest path does not necessarily have most cache misses
 - e.g., second-longest path could have more cache misses and in fact take longer

WCET with Caches – What can we do?

- 1 Just ignore them
 - **no:** This will lead to huge over- or underestimations up to factor of 100 on modern processors.
- 2 Find longest path with ILP technique, and then simulate it to get cache misses/hits
 - **no:** Longest path does not necessarily have most cache misses
 - e.g., second-longest path could have more cache misses and in fact take longer
- 3 Find all possible paths, and then simulate them to get cache misses/hits

WCET with Caches – What can we do?

- 1 Just ignore them
 - **no**: This will lead to huge over- or underestimations up to factor of 100 on modern processors.
- 2 Find longest path with ILP technique, and then simulate it to get cache misses/hits
 - **no**: Longest path does not necessarily have most cache misses
 - e.g., second-longest path could have more cache misses and in fact take longer
- 3 Find all possible paths, and then simulate them to get cache misses/hits
 - **no**: computationally infeasible, too many paths (see lecture 6)

WCET with Caches – What can we do?

- 1 Just ignore them
 - **no**: This will lead to huge over- or underestimations up to factor of 100 on modern processors.
- 2 Find longest path with ILP technique, and then simulate it to get cache misses/hits
 - **no**: Longest path does not necessarily have most cache misses
 - e.g., second-longest path could have more cache misses and in fact take longer
- 3 Find all possible paths, and then simulate them to get cache misses/hits
 - **no**: computationally infeasible, too many paths (see lecture 6)
- 4 Extend ILP: provide constraints for min/max number of hits/misses

WCET with Caches – What can we do?

- 1 Just ignore them
 - **no:** This will lead to huge over- or underestimations up to factor of 100 on modern processors.
- 2 Find longest path with ILP technique, and then simulate it to get cache misses/hits
 - **no:** Longest path does not necessarily have most cache misses
 - e.g., second-longest path could have more cache misses and in fact take longer
- 3 Find all possible paths, and then simulate them to get cache misses/hits
 - **no:** computationally infeasible, too many paths (see lecture 6)
- 4 Extend ILP: provide constraints for min/max number of hits/misses
 - **advantage:** tight estimate

WCET with Caches – What can we do?

Use **two variables per instruction** (hit/miss count for each instruction: x_i^{hit}, x_i^{miss}) \rightarrow **might get computationally infeasible**

- New ILP Problem:

$$\text{WCET} := \max_{x^{hit}, x^{miss}} \sum_{i=1 \dots N} c_i^{hit} x_i^{hit} + c_i^{miss} x_i^{miss} \quad (1)$$

- very inefficient
- practically, this only works for small programs, not in industry

WCET with Caches – L-Blocks

Solution:

WCET with Caches – L-Blocks

Solution:

Group instructions to blocks again

WCET with Caches – L-Blocks

Solution:

Group instructions to blocks again - but how?

WCET with Caches – L-Blocks

Solution:

Group instructions to blocks again - but how?

Observation:

Only two kinds of instructions can have a cache miss:

- instructions that map to the beginning of a cache line
- instructions that are the first one in a BB

All other instructions will have a cache hit.

WCET with Caches – L-Blocks

Solution:

Group instructions to blocks again - but how?

Observation:

Only two kinds of instructions can have a cache miss:

- instructions that map to the beginning of a cache line
- instructions that are the first one in a BB

All other instructions will have a cache hit.

- If we group all instructions that are in the same BB and same cache line (called *line-blocks* or *l-blocks*), each group has only two possible execution times:
 - 1 first instruction = miss & all subsequent = hit
 - 2 first instruction = hit & all subsequent = hit

WCET with Caches – L-Blocks

Alternative definition:

l-blocks = max. sequence of code within a BB, s.t. when 1st instruction of it is accessed, either entire *l-block* is in cache or none

WCET with Caches – L-Blocks

Alternative definition:

l-blocks = max. sequence of code within a BB, s.t. when 1st instruction of it is accessed, either entire *l*-block is in cache or none
re-formulate the ILP-problem with *l*-blocks

$$\text{WCET} := \sum_{i=1}^N \sum_{j=1}^{n_i} c_{i,j}^{\text{hit}} x_{i,j}^{\text{hit}} + c_{i,j}^{\text{miss}} x_{i,j}^{\text{miss}}, \quad x_i = x_{i,j}^{\text{hit}} + x_{i,j}^{\text{miss}} \quad (2)$$

Now, *i* sums over BBs again, and *j* over the *l*-blocks. More efficient than *per instruction*.

WCET with Caches – L-Blocks

Alternative definition:

l-blocks = max. sequence of code within a BB, s.t. when 1st instruction of it is accessed, either entire *l*-block is in cache or none
re-formulate the ILP-problem with *l*-blocks

$$\text{WCET} := \sum_{i=1}^N \sum_{j=1}^{n_i} c_{i,j}^{\text{hit}} x_{i,j}^{\text{hit}} + c_{i,j}^{\text{miss}} x_{i,j}^{\text{miss}}, \quad x_i = x_{i,j}^{\text{hit}} + x_{i,j}^{\text{miss}} \quad (2)$$

Now, *i* sums over BBs again, and *j* over the *l*-blocks. More efficient than *per instruction*.

But: How do we know if an access will be a cache hit or miss?

WCET with Caches – Cache Conflicts

so far: we only looked at caches, but not how they depend on control flow

- **remember:** cache is small. Often the program does not fit in there.
- this means, sometimes data in the cache has to be replaced (“evicted”)
- evictions may result in cache misses
- the *control flow* (CFG) defines the evictions
- we need to link the CFG to the objective function seen before, i.e., we need further constraints for our ILP (What if we don't provide them?)

WCET with Caches – Cache Conflicts

The *Cach-Conflict-Graph* connects control flow with the cache effects

- **example:** see lecture
- essentially we construct a smaller CFG again, with nodes and edges
- resulting equations = more constraints for our ILP

Finally: Extended ILP objective function + structural constraints from CFG + more constraints from CCG = our new, **cache-aware WCET analysis**